



AutoFL: Enabling Heterogeneity-Aware Energy Efficient Federated Learning

Young Geun Kim

Soongsil University
younggeun.kim@ssu.ac.kr

Carole-Jean Wu

Arizona State University
Facebook AI Research
carole-jean.wu@asu.edu

ABSTRACT

Federated learning enables a cluster of decentralized mobile devices at the edge to collaboratively train a shared machine learning model, while keeping all the raw training samples on device. This decentralized training approach is demonstrated as a practical solution to mitigate the risk of privacy leakage. However, enabling efficient FL deployment at the edge is challenging because of non-IID training data distribution, wide system heterogeneity and stochastic-varying runtime effects in the field. This paper jointly optimizes time-to-convergence and energy efficiency of state-of-the-art FL use cases by taking into account the stochastic nature of edge execution. We propose AutoFL by tailor-designing a reinforcement learning algorithm that learns and determines which K participant devices and per-device *execution targets* for each FL model aggregation round in the presence of stochastic runtime variance, system and data heterogeneity. By considering the unique characteristics of FL edge deployment judiciously, AutoFL achieves 3.6 times faster model convergence time and 4.7 and 5.2 times higher energy efficiency for local clients and globally over the cluster of K participants, respectively.

CCS CONCEPTS

• **Computer systems organization** → *Embedded systems; Cloud computing*; • **Computing methodologies** → *Machine learning*.

KEYWORDS

Federate learning, mobile devices, heterogeneity, energy efficiency, reinforcement learning

ACM Reference Format:

Young Geun Kim and Carole-Jean Wu. 2021. AutoFL: Enabling Heterogeneity-Aware Energy Efficient Federated Learning. In *MICRO'21: IEEE/ACM International Symposium on Microarchitecture, October 18-22, 2021, Athens, Greece*. ACM, New York, NY, USA, 16 pages. <https://doi.org/10.1145/3466752.3480129>

1 INTRODUCTION

The ever increasing computational capacities and efficiencies of smartphones have enabled a large variety of machine learning (ML)

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MICRO'21, October 18–22, 2021, Athens, Greece

© 2021 Association for Computing Machinery.

ACM ISBN 978-1-4503-8557-2/21/10...\$15.00

<https://doi.org/10.1145/3466752.3480129>




Figure 1: The performance-per-watt (PPW) energy efficiency of FL execution can be significantly improved by up to 5.4x with judicious selections of participant devices and execution targets (Performance and O_{FL} — Section 5 for methodology details).

applications at the edge [128], such as image recognition [36], virtual assistant [4, 8], language translation [38], automatic speech recognition [39], and recommendation [53]. As the mobile ML system stack matures [6, 16, 92, 98, 100, 108, 117], on-device inference becomes more efficient with innovations in algorithmic optimizations [45, 67, 82, 109, 115, 126, 145], neural network architecture optimizations [46, 109, 114, 127], and the availability of programmable accelerators [7, 37, 50, 51, 100, 105, 106]. While on-device inference is becoming more ubiquitous [14, 29, 41, 58, 63, 66, 103, 104, 122, 128, 143], performing ML model training in the cloud remains the standard practice for most use cases [1, 29, 44, 55, 83, 90, 103] due to the substantial computation and memory requirements [9, 35, 70, 102, 124, 125, 134, 142].

Recently, federated learning (FL) enables smartphones to collaboratively train a shared ML model while keeping all the raw data on device [11, 34, 52, 64, 71, 77, 84, 116, 120, 135, 137]. This decentralized approach is a practical way to mitigate the risk of privacy leakage when training deep neural networks (DNNs), as only the model gradients—not individual data samples—go to the cloud to update the shared model [12, 43, 73]. The shared model is trained iteratively using the model gradients from a large collection of participating smartphones. While FL has shown great promise for privacy sensitive tasks including sentiment learning, next word prediction, health monitoring, and item ranking [11, 43, 71], its deployment is still in a nascent stage.

A common practice for enabling efficient FL deployment at the edge is to maximize the computation-communication ratio by employing fewer participant devices with higher per-device training iterations [73, 84, 112]. In particular, FedAvg has been considered as the de facto FL algorithm [64, 84]. For each aggregation round, FedAvg trains a model over E epochs using stochastic gradient descent (SGD) with minibatch size of B on K selected devices, where K

is a small subset of N devices participating in the FL. The K devices then upload the respective model gradients to the cloud, where the gradients get averaged into the shared model. By allowing lower K , FedAvg reduces the amount of data transmission for each aggregation round. Various previous works have been also proposed to improve the accuracy of trained models [28, 71, 74] or security robustness [34, 76, 79] atop the FedAvg algorithm.

While these advancements open up the possibility of efficient FL deployment, a fundamental challenge remains—deciding which K devices to participate in each aggregation round for a given (B, E, K) ¹, and selecting the *execution target* for model training on a participating device. State-of-the-art approaches randomly select K participants from a total of N devices [11, 64, 73, 84, 112], leaving a significant energy efficiency and model convergence co-optimization opportunity on the table (Figure 1).

System heterogeneity and stochastic runtime variance: At the edge, there are over two thousand unique systems-on-a-chip (SoCs) with different compute resources, including CPUs, graphics processing units (GPUs), and digital signal processors (DSPs), in more than ten thousand different smart devices [63, 128]. The high degree of system heterogeneity introduces varying, potentially large, performance gaps across smartphones participating in FL. Furthermore, mobile execution is stochastic by nature [32, 33, 128]. The performance variability can stem from interference between and within applications [129] as well as the stability of network. Altogether, these factors lead to the straggler problem—training time of each aggregation round is limited by the slowest participating smartphone. To mitigate the straggler problem, several previous works built atop FedAvg by excluding stragglers from each round [84] or allowing partial updates from the straggler [73]. However, these approaches sacrifice accuracy.

Data heterogeneity: Varying characteristics of training data per participating device introduce additional challenge to efficient FL execution [15, 75]. To guarantee model convergence, it is important to ensure that training data is independently and identically distributed (IID) across the participating devices [15, 40, 131]. For example, if a model classifies images into 10 distinct label categories, the data samples are IID if each individual device has independent data representing all 10 categories [120]. However, in a realistic environment, the training data samples on each device are usually based on the user behavior, preference, or both. Thus, local training samples for any particular user will be unrepresentative of the population, deferring convergence [75, 84]. To mitigate data heterogeneity, previous approaches excluded the non-IID devices [17, 18], used a warm up model [141], or shared data across a subset of the participant devices [28, 71]. However, none has considered both data and system heterogeneity with runtime variance.

Furthermore, there has been little work on energy efficiency optimization for FL. Most prior work assumed FL is only activated when smartphones are plugged into wall power, due to the significant energy consumption of model training [17, 99, 120, 130, 135]. Unfortunately, this assumption has limited the practicality of FL,

increasing model convergence time and degrading model accuracy [77, 120]. Energy efficient FL could enable on-device training anytime, with better model quality and user experience.

To tackle challenges in a realistic execution environment, this paper proposes a learning-based energy optimization framework—*AutoFL*—that selects K participants as well as execution targets to guarantee model quality, while maximizing the energy efficiency of individual participants (or the cluster of all participants), for FL. The optimization is performed by considering the presence of system and data heterogeneity and runtime variance. Since the optimal decision varies with NN characteristics, FL global parameters, profiles of participating devices, distributions of local training samples, and stochastic runtime variance, the design space is massive and infeasible to enumerate. Thus, we design a reinforcement-learning technique. For each aggregation round, AutoFL observes the NN characteristics, FL global parameters, and system profiles of devices (including interference intensity, network stability, and data distributions). It then selects the participant devices for the round and simultaneously determines the execution target for each participant, to maximize energy efficiency while satisfying the training accuracy requirements. The result of the decision is measured and fed back to AutoFL, allowing it to continuously learn and predict the near-optimal action for subsequent rounds.

AutoFL is implemented and runs on the centralized, model aggregation server. We evaluate our design using 200 mobile systems comprising three major categories: high, medium, and low performance. The key contributions of this work are as follows:

- We present an in-depth performance and energy efficiency characterization for FL by considering a realistic edge-cloud execution environment. The results show that the optimal participant selection and resource allocation in FL can vary significantly with NN characteristics, the varying degree of data and system heterogeneity, and the stochastic nature of mobile execution (Section 3).
- We propose a FL energy optimization framework, AutoFL, that identifies near-optimal participant selection and resource allocation at runtime, enabling heterogeneity-aware energy efficient federated learning (Section 4).
- To demonstrate the feasibility and practicality, we design, implement, and evaluate AutoFL for a variety of FL use cases in the edge-cloud environment (Section 6). Real-system experiments show that AutoFL improves energy efficiency of individual participant devices as well as for the cluster of all participating devices by an average of 4.7x and 5.2x, respectively, while also satisfying the accuracy requirement. By considering runtime variance along with system and data heterogeneity, AutoFL achieves an average of 49.8% and 39.3% higher energy efficiency, compared to the state-of-the-art techniques, FedNova [121] and FEDL [26], respectively.

2 BACKGROUND

2.1 Federated Learning

To improve data privacy for ML training, federated learning (FL) is introduced by allowing devices at the edge, such as smartphones, to collaboratively train a shared ML model while keeping all user data locally on the device [11, 34, 52, 64, 71, 77, 84, 116, 120, 135, 137].

¹The FL global parameters B , E , and K are usually determined by the service providers on the basis of service-level accuracy requirements as well as the computation and memory capabilities of edge devices [11, 64].




Figure 2: Overview for federated learning (FL).

Figure 2 depicts the overall system architecture for the FL baseline [11, 64, 84]. There are two entities in the FL system: an *aggregation server* as the model owner and a collection of local devices (data owners). Given N local devices, the server first initializes a global deep learning model and its global parameters by specifying the number of local training epochs E , the local training minibatch size B , and the number of participant devices K . (B, E, K) is determined by the FL-based services [11, 84].

In each aggregation round, the server selects K participants among the N devices (Step ①) and broadcasts the global model to the selected devices (Step ②). Each participant independently trains the model using local data samples with the batch size of B over E epochs (Step ③). Once the local training step is complete, the computed model gradients are sent back to the server (Step ④). The server then aggregates the local gradients and calculates their average [84] to update the global model (Step ⑤). The steps are repeated until a desirable accuracy is achieved.

2.2 Realistic Execution Environment

System heterogeneity, runtime variance, and data heterogeneity form a massive optimization space for FL. Figure 3 illustrates FL execution in a realistic environment. In this example, a cluster of two hundred devices participate in FL. Depending on the performance of an individual device (i.e., high-end, mid-end, or low-end smartphone) and the availability of co-processors, such as GPUs, DSPs, or neural processing units (NPUs), the training time varies. Extensive system heterogeneity introduces large performance gaps among the devices, leading to the straggler problem [73, 77, 84, 120, 135].

In addition, stochastic runtime variance can exacerbate the straggler problem. Depending on the amount of on-device interference and the execution conditions, such as ambient temperature and network signal strength, the execution time performance of each participant—the training time per round (*Computation Time*) and the model gradient aggregation time (*Communication Time*)—is highly dynamic. Finally, not all participant devices possess IID training samples. Heterogeneous data across devices can significantly deteriorate FL model convergence and quality.

3 MOTIVATION

This section presents FL system characterization results. We examine the design space covering three important axes: energy efficiency, convergence time, and accuracy.




Figure 3: Optimization space of FL is large due to the scale of decentralized training, system heterogeneity, data heterogeneity, and runtime variance.

3.1 Impact of FL Global Parameters and NN Characteristics

The optimal cluster of participating devices depends on the FL global parameters and the resource requirement of NN models. From the system's perspective, the global parameters determine the amount of computations performed on each individual device. Figure 4 compares the energy efficiency achieved under the four different FL global parameter settings (S1 to S4 defined in Table 5 of Section 5.2) for training the CNN model with the MNIST dataset (CNN-MNIST) over the eight different combinations of participant devices (C0 to C7 defined in Table 4). The optimal device cluster changes from C1 to C2, C3 and C4 when the global parameter setting changes from S1 to S2, S3, and S4, respectively.

When the number of computations assigned to each device is large (i.e., S1), including more number of high-end devices is beneficial as they exhibit 1.7x and 2.5x better training time, compared to mid-end and low-end devices, respectively, due to powerful CPUs and co-processors along with larger size of cache and memory. On the other hand, when the number of computations assigned to each device decreases (i.e., from S1 to S2 and S3), including mid-end and low-end devices along with the high-end devices results in better energy efficiency since their lower power consumption (35.7% and 46.4% compared to high-end devices respectively) amortizes the performance gap. If K is decreased (i.e., from S3 to S4), reducing the number of high-end devices is beneficial as the devices can stay idle during the round — though mid-end devices exhibit longer




Figure 4: Depending on FL global parameters and NN model resource needs, the optimal clusters of K participating devices are C1, C2, C3, and C4 for the four different global parameter settings, respectively. The detailed descriptions for the settings and the clusters are in Table 5 and Table 4 of Section 5, respectively. Striped bars indicate the optimal cluster.




Figure 5: With runtime variance from various sources, the optimal cluster of K participating devices shifts from C3 to C1 and C5. PPW is normalized to C0 with no runtime variance.

training-time-per-round than high-end devices, similar to S3, the mid-end devices have better energy efficiency in this case.

When we use the LSTM model with Shakespeare dataset (denoted as LSTM-Shakespeare), the optimal device cluster over S1–S4 is C3, C4, C5, and C5, respectively, compared with CNN-MNIST’s C1, C2, C3, C4. In the case of CNN-MNIST, high-end devices with more powerful mobile SoCs exhibit better performance and energy efficiency than mid- and low-end devices, due to the compute-intensive CONV and FC layers. On the other hand, for LSTM-Shakespeare, the energy efficiency of mid- and low-end devices is similar to that of high-end devices. The reason is that the performance variation among the devices diminishes (from 2.1x to 1.5x) due to the memory operations, so that the low power consumption of the mid- and low-end devices compensates for their performance loss.

3.2 Impact of Runtime Variance

The optimal cluster of participants also significantly varies along with the runtime variance. Figure 5(a) compares the energy efficiency when on-device interference is absent and when the network signal is stable. In such an ideal execution environment, the most energy efficient cluster is C3, balancing the trade-off between the training-time-per-round and power consumption of different device categories — C3 achieves 3.2x higher energy efficiency than the baseline C0. In the presence of on-device interference, the optimal cluster becomes C1 (Figure 5(b)), whereas when the network signal strength is weak, the optimal cluster switches to C5 (Figure 5(c)).




Figure 6: (a) FL Model quality and (b) energy efficiency change with varying levels of data heterogeneity.

Intuitively, in the presence of on-device interference, it is more beneficial to select high-end devices to participate in FL— they have a high computation and memory capabilities [63], achieving 2.0x and 3.1x better performance compared to mid-end and low-end devices, respectively. On the other hand, when the network signal strength is poor, the communication time and energy on each device is significantly increased [25, 62] (4.3x, on average). In this case, the impact of performance gap among different categories of devices decreases along with the decreased portion of computation time. For this reason, including low-power devices is beneficial in terms of energy efficiency due to lower computation and communication power consumption.

3.3 Impact of Data Heterogeneity

Participant device selection strategies that ignore data heterogeneity lead to sub-optimal FL execution. Figure 6(a) depicts the convergence patterns for CNN-MNIST over varying degrees of data heterogeneity—the x-axis shows the consecutive FL rounds and the y-axis shows the model accuracy. Here, Non-IID (M%) means M% of K participant devices have non-IID data where a portion of the samples of each data class is distributed following Dirichlet distribution with a concentration parameter of 0.1 [15, 57, 72, 75, 78], while the rest of devices have all the data classes independently — the smaller concentration parameter, the more each data class is concentrated on one device.

Data heterogeneity can significantly affect model convergence—when devices with non-IID data participate in FL, the convergence time is significantly increased compared to the ideal IID scenario. The increased convergence time eventually deteriorates FL energy efficiency. Figure 6(b) illustrates the large (>85%) energy efficiency gap between the ideal device selection scenario and the sub-optimal selection scenarios with non-IID data.

4 AUTOFL

To capture stochastic runtime variance in the presence of system and data heterogeneity, we propose an adaptive prediction mechanism based on reinforcement learning², called AutoFL. In general,

²We exploit RL instead of other statistical methods, such as Gaussian Process, since RL has the following advantages: (1) faster training and inference due to lower complexity [10, 93], (2) higher sample efficiency (i.e., the amount of experiments to reach a certain level of accuracy) [59, 86, 133], and (3) higher prediction accuracy under stochastic variance [63, 93].

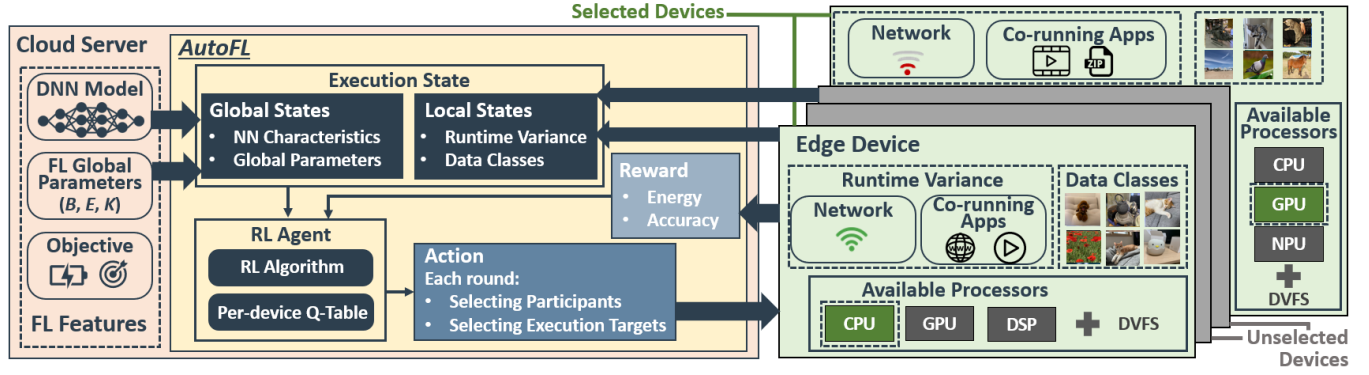


Figure 7: AutoFL Design Overview.

an RL agent learns a policy to select the best action for a given state through accumulated rewards [93]. In the context of FL, given an NN and the corresponding global parameters, AutoFL learns to select a near-optimal cluster of participants and an energy-efficient execution target for each one in every aggregation round.

Figure 7 shows the AutoFL design overview. In each FL aggregation round, AutoFL observes the global configurations of FL, including the target NN and the global parameters. In addition, it collects the execution states of participant devices, including their resource usage and network stability³, and the number of data classes each device has. Based on the information, AutoFL identifies the current execution state⁴. For the identified state, AutoFL selects participant devices that are expected to maximize the energy efficiency of FL, while satisfying the accuracy requirement. It also determines the execution target for each device to additionally improve the local energy efficiency. The selections are based on per-device lookup tables (i.e., Q-tables) that contain the accumulated rewards of previous selections. After the gradient updates are aggregated in the server, AutoFL measures the results (i.e., training time, energy consumption, and test accuracy) to calculate the reward—how the selected action improves global as well as local energy efficiency and accuracy. Finally, it updates the per-device Q-table with the calculated reward. Optimizing a system through RL involves three important design requirements.

High prediction accuracy: High prediction accuracy is essential to the success of an RL-based approach. To handle the dynamic execution environment of FL, it is important to model the core components—state, action, and reward—in a realistic environment. We define these components in accordance with our observations (Section 4.1). In addition to the core components, avoiding local optima is also important. The fine balance between exploitation versus exploration is at the heart of RL [30, 65]. If an RL agent always exploits an action with the temporary highest reward, it can get

stuck in a local optimum. On the other hand, if it keeps exploring all possible actions, convergence of RL may take too long. To tackle this challenge, we employ the epsilon-greedy algorithm. This algorithm is one of the commonly-adopted algorithms [63, 81, 91, 93] due to its effectiveness and simplicity (Section 4.2) — it achieves prediction accuracy similar to that of other complex algorithms, such as Exp3, Softmax, UCB, and Thompson Sampling, but incurs less overhead [21, 118].

Low training and inference overhead: To minimize the timing and energy overhead of on-device RL, AutoFL expedites the RL training by enabling devices within the same performance category to share the learned results — in a realistic environment, each user can experience different degree of the data heterogeneity and runtime variance, and thus sharing the learned results across the devices complements one another. We present the training time reduction of this approach in Section 6.4.

The inference latency of per-device RL models determines the decision making performance of AutoFL. Thus, among the various RL implementation choices, e.g., Q-learning [20] and deep RL [85], Q-learning is most suitable to this work — it achieves low training and inference latency using look-up tables, whereas deep RL usually suffers longer latency because of forward and backward propagation of DNNs [93].

Scalability: As energy efficient FL can enable many more devices to participate in FL, scalability to a large number of devices is crucial. To permit a large number of participant, AutoFL can exploit a shared Q-table for devices within the same performance category — additional clustering algorithm can be used along with the AutoFL for binding similar category of devices. By updating the shared Q-table instead of all the per-device Q tables, AutoFL can handle the large number of devices, at the expense of a small prediction accuracy loss (see details in Section 6.4).

4.1 AutoFL Reinforcement Learning Design

We define the core RL components—*State*, *Action*, and *Reward*—to formulate the optimization space for AutoFL.

State: Based on the observations presented in Section 3, we identify states that are critical to energy-efficient FL execution. Table 1 summarizes the states.

³AutoFL relies on the resource usage and network bandwidth information collected by the de facto FL protocol [11] — the protocol collects such information to ensure robust model training. Note, to keep system usage information private, it is possible to run training of per-device Q-table locally, without sharing the information with the cloud server at the expense of increased training cost (336.2 μ s in low-end device including communication cost.).

⁴AutoFL assumes that runtime variance does not significantly vary during a short time period (e.g., in the order of 10-100 milliseconds).

Table 1: State features for AutoFL.

State		Description	Discrete Values
NN-related Features	S_{CONV}	# of CONV layers	Small (<10), medium (<20), large (<30), larger (>=40)
	S_{FC}	# of FC layers	Small (<10), large (>=10)
	S_{RC}	# of RC layers	Small (<5), medium (<10), large (>=10)
Global Parameters	S_B	Batch size	Small (<8), medium (<32), large (>=32)
	S_E	# of local epochs	Small (<5), medium (<10), large (>=10)
	S_K	# of participant devices	Small (<10), medium (<50), large (>=50)
Runtime Variance	S_{Co_CPU}	CPU utilization of co-running apps	None (0%), small (<25%), medium (<75%), large (<=100%)
	S_{Co_MEM}	Memory usage of co-running apps	None (0%), small (<25%), medium (<75%), large (<=100%)
	$S_{Network}$	Network bandwidth	Regular (>40Mbps), bad (<=40Mbps)
Data Classes	S_{Data}	# of data classes for this round	Small (<25%), medium (<100%), large (=100%)

First, the energy efficiency of devices highly depends on NNs and the given global parameters. To model the impact of NN characteristics and global parameters, we identify states with layer types that are deeply correlated with the energy efficiency and performance of on-device training. We test the correlation between each layer type and the energy efficiency by calculating the squared correlation coefficient (ρ^2) [144]. We find convolution layers (CONV), fully-connected layers (FC), and recurrent layers (RC) impact energy efficiency differently due to their respective compute- and/or memory-intensive natures. Thus, we identify S_{CONV} , S_{FC} , and S_{RC} to represent the number of CONV, FC, and RC layers in a NN, respectively. We also identify S_B , S_E , and S_K as the global parameters of batch size, the number of local epochs, and the number of participant devices, respectively.

The energy efficiency of participating devices is highly subject to the runtime variance—namely, on-device interference and network instability. To model on-device interference, we identify the per-device states of S_{Co_CPU} and S_{Co_MEM} to represent CPU utilization and memory usage of co-running applications, respectively. We also model per-device network instability with $S_{Network}$ to represent the network bandwidth of the respective wireless network (e.g., Wi-Fi, LTE, and 5G). In addition, data heterogeneity also has a substantial impact on the FL convergence time and energy efficiency. Therefore, to model the impact of data heterogeneity on the FL efficiency, we identify S_{Data} which stands for the number of data classes that each device has for an aggregation round⁵.

When a feature has a continuous value, it is difficult to define the state in a discrete manner for the lookup table of Q-learning [20, 63, 91]. To convert the continuous features into discrete values, we applied the DBSCAN clustering algorithm to each feature [20, 63]—DBSCAN determines the optimal number of clusters for the given data. The last column of Table 1 summarizes the discrete values.

Action: Actions in reinforcement learning represent the tunable control knobs of a system. In FL, we define the actions in two levels. At the global level, we define the selection of participant devices as an action. For each selected device, we define the selection of on-device execution targets available for training execution, such as CPUs, GPUs, or DSP, as another action. The execution targets

⁵Note, while per-device data class distribution is also used in a number of prior works [19, 42, 80, 87, 132, 140], it reveals additional information. To mitigate privacy concerns, AutoFL can exclude the use of data class distribution from the RL state encoding. In this case, although the convergence of RL can become slightly slower (from 50-80 to 70-90 aggregation rounds to convergence), AutoFL can still select the near-optimal participant cluster and achieves an average of 90.8% prediction accuracy.

can be augmented to include CPU dynamic voltage and frequency scaling (DVFS) settings, to exploit the performance slack caused by stragglers for further energy saving.

Reward: In RL, rewards track the optimization objective of the system. To represent the main optimization axes, we encode three rewards: R_{energy_local} , R_{energy_global} , and $R_{accuracy}$. R_{energy_local} is the estimated energy consumption of each individual device and R_{energy_global} is the estimated energy consumption of all participating devices. $R_{accuracy}$ represents the test accuracy of NN.

We estimate R_{energy_local} and R_{energy_global} as follows. For each selected participant device, we first calculate the computation energy, E_{comp} . When the CPU is selected as the execution target, E_{comp} is calculated using a utilization-based CPU power model [13, 54, 63, 138], as in (1), where E_{core}^i is the power consumed by the i th core, t_{busy}^f and t_{idle} are the time spent in the busy state at frequency f and that in the idle state, respectively, and P_{busy}^f and P_{idle} are the power consumed during t_{busy}^f at f and that during t_{idle} , respectively.

$$E_{comp} = \sum_i E_{core}^i, \quad (1)$$

$$E_{core} = \sum_f (P_{busy}^f \times t_{busy}^f) + P_{idle} \times t_{idle}$$

Similarly, if GPU is selected as the execution target, E_{comp} is calculated using the GPU power model [24] as in (2). Note that t_{busy}^f and t_{idle} for CPU/GPU are obtained from *procf*s and *sysfs* in the Linux kernel [20], whereas P_{busy}^f and P_{idle} for CPU/GPU are obtained by power measurement of the CPU/GPU at each frequency in the busy and idle states, respectively⁶. Those values are obtained for representative edge device categories (i.e., high-end, mid-end, and low-end devices) and stored in a look-up table of AutoFL.

$$E_{comp} = \sum_f (P_{busy}^f \times t_{busy}^f) + P_{idle} \times t_{idle} \quad (2)$$

After calculating the computation energy, we calculate the communication energy, E_{comm} , for each selected participant using the signal strength-based energy model [62], as in (3). Here, t_{TX} is the latency measured while transmitting the gradient updates, and P_{TX}^S

⁶Although we only present the energy estimation for CPU and GPU in this paper due to the limited programmability of on-device training, a similar practice can also be used for other co-processors, such as DSPs and NPUs [41, 123].

is the power consumed by the device during t_{TX} at signal strength S . Note P_{TX}^S is obtained by measuring the power consumption of devices at each signal strength when transmitting data.

$$E_{comm} = P_{TX}^S \times t_{TX} \quad (3)$$

We also calculate the idle energy, E_{idle} , for non-selected devices, as in (4). Here, t_{round} is the time spent during the training round.

$$E_{idle} = P_{idle} \times t_{round} \quad (4)$$

Based on the estimated energy values, R_{energy_local} is calculated for each device, as in (5), where S_t represents a set of selected participants.

$$\begin{aligned} & \text{if } device \in S_t \\ & \quad R_{energy_local} = E_{comp} + E_{comm} \\ & \text{else} \\ & \quad R_{energy_local} = E_{idle} \end{aligned} \quad (5)$$

In addition, R_{energy_global} is calculated for a cluster of all N participating devices, as in (6), based on the R_{energy_local} .

$$R_{energy_global} = \sum_i^N R_{energy_local} \quad (6)$$

Since the energy estimation is based on the measured latency, its mean absolute percentage error is 7.3%—low enough to identify the optimal participants and execution targets.

To ensure AutoFL selects participants and their corresponding execution targets that maximize energy efficiency while satisfying the accuracy requirements, the reward R is calculated as in (7)⁷, where $R_{accuracy_prev}$ is the test accuracy of the training NN from the previous round. α and β are, respectively, the weights for the accuracy and the amount of accuracy improvement, which is directly related to the convergence speed.

$$\begin{aligned} & \text{if } R_{accuracy} - R_{accuracy_prev} \leq 0, \\ & \quad R = R_{accuracy} - 100 \\ & \text{else} \\ & \quad R = -R_{energy_global} - R_{energy_local} \\ & \quad \quad + \alpha R_{accuracy} + \beta (R_{accuracy} - R_{accuracy_prev}) \end{aligned} \quad (7)$$

If the selected action fails to improve the accuracy from the previous round, the reward is $R_{accuracy} - 100$ (i.e., how much the accuracy is far from 100%) to avoid choosing the action for the next inference. Otherwise, the reward is calculated for each device based on the global energy, local energy, accuracy, and the amount of accuracy improvement.

4.2 AutoFL Implementation Detail

AutoFL is built based on Q-learning. To strike a balance between exploitation and exploration in RL, it employs the epsilon-greedy algorithm with a uniformly random action, based on a pre-specified exploration probability. For the rest, AutoFL chooses an action with the highest reward.

⁷We include the energy consumption as a reward to model the impact of selections on the global and local energy efficiency. We include accuracy to model the impact of selections on model quality.

Algorithm 1 Training the Q-learning model

Variable: S_{global}, S_{local}, A

S_{global} is the global state

S_{local} is the local state

A is the action (execution target)

Constants: γ, μ, ϵ

γ is the learning rate

μ is the discount factor

ϵ is the exploration probability

Initialize $Q(S_{global}, S_{local}, A)$ as random values

Repeat (whenever an aggregation round begins):

Observe global state and store in S_{global}

Observe local state for each device and store in S_{local}

if $\text{rand}() < \epsilon$ **then**

Choose K participants randomly

Choose action A randomly for selected participants

else

Sort devices by $Q(S_{global}, S_{local}, A)$

Choose at most top K participants

Choose action A with the largest $Q(S_{global}, S_{local}, A)$

Run training on a target defined by A in each device

(when local training and aggregation end)

Estimate $R_{energy_global}, R_{energy_local}$, and obtain

$R_{accuracy}$

Calculate reward R

Observe new global state S'_{global}

Observe new local state S'_{local}

Sort devices by $Q(S'_{global}, S'_{local}, A')$

Choose at most top K participants

Choose action A' with the largest $Q(S'_{global}, S'_{local}, A')$

$Q(S_{global}, S_{local}, A) \leftarrow Q(S_{global}, S_{local}, A) + \gamma[R + \mu Q(S'_{global}, S'_{local}, A') - Q(S_{global}, S_{local}, A)]$

$S \leftarrow S'$

In Q-learning, the value function $Q(S_{global}, S_{local}, A)$ takes the global state S_{global} , the local state S_{local} , and the action A as parameters in the form of a lookup table (Q-table). Algorithm 1 shows the detailed algorithm for training the per-device Q-table. At the beginning, AutoFL initializes the Q-tables with random values. At runtime, it observes S_{global} and S_{local} for each aggregation round by checking the NN characteristics, runtime variance, and data heterogeneity. It evaluates a random value compared with ϵ ⁸. If the random value is smaller than ϵ , AutoFL selects participants randomly and determines A for exploration. Otherwise, it sorts the devices by $Q(S_{global}, S_{local}, A)$ and selects the top K devices.

Next, AutoFL chooses A with the largest $Q(S_{global}, S_{local}, A)$ for each selected participant. After the local training and the aggregation end, AutoFL estimates R_{energy_local} and R_{energy_global} as explained in Section 4.1. In addition, it obtains $R_{accuracy}$ and $R_{accuracy_prev}$. Based on these values, it calculates the reward R as in (7) of Section 4.1. Afterward, AutoFL observes the new states and

⁸Note that we use 0.1 for ϵ based on our sensitivity analysis.

Table 2: Amazon EC2 instance specification.

Level	Instance	Performance (GFLOPS)	RAM (GB)
H	m4.large	153.6	8
M	t3a.medium	80.0	4
L	t2.small	52.8	2

Table 3: Mobile device specification.

Device	CPU	GPU
Mi8Pro (H)	Cortex-A75 (2.8GHz) 23 V/F steps 5.5 W	Adreno 630 (0.7GHz) 7 V/F steps 2.8 W
Galaxy S10e (M)	Mongoose (2.7GHz) 21 V/F steps 5.6 W	Mali-G76 (0.7GHz) 9 V/F steps 2.4 W
Moto X Force (L)	Cortex-A57 (1.9GHz) 15 V/F steps 3.6 W	Adreno 430 (0.6GHz) 6 V/F steps 2.0 W

chooses the corresponding participants and execution targets using $Q(S'_{global}, S'_{local}, A')$. It then updates the $Q(S_{global}, S_{local}, A)$ based on the equation in Algorithm 1. In the equation, γ and μ are hyperparameters that represent the learning rate and discount factor, respectively. We set γ and μ in accordance with the sensitivity evaluation. Section 5.3 describes hyperparameter tuning.

After learning is completed—i.e., the largest $Q(S_{global}, S_{local}, A)$ value is converged for each S_{global} and S_{local} —AutoFL uses the per-device Q -tables to select participants and the corresponding A , which maximizes $Q(S_{global}, S_{local}, A)$ for the observed S_{global} and S_{local} . Note, among the devices with the same Q value, AutoFL randomly selects participants to avoid biased selection [73, 74].

5 EXPERIMENTAL METHODOLOGY

5.1 System Measurement Infrastructure

We set up an edge-cloud FL system that consists of 200 mobile devices ($N = 200$) and one model aggregation server. Similar FL system infrastructures have been used in prior works [28, 64, 73, 84]. We emulate the performance of FL execution by using Amazon EC2 instances [5] that provide the same theoretical GFLOP performance as the three representative smartphone categories: high-end (H), mid-end (M), and low-end (L) devices. Table 2 summarizes the system profiles. Among the 200 instances, there are 30 H, 70 M, and 100 L devices, representative of in-the-field system performance distribution [128].

For model aggregation, we connect the aforementioned systems to a high-performance Amazon EC2 system instance, c5d.24xlarge, which has a theoretical performance of 448 GFLOPS and is equipped with 32GB of RAM. We perform power measurement directly using an external Monsoon Power Meter [88] for the three smartphones during on-device training (implemented with DL4j [27]): Mi8Pro [50], Galaxy S10e [107], and Moto X Force [89] (Table 3). Similar power measurement methodologies are used in prior works [95, 96, 111].

Table 4: Cluster of devices used for characterization.

Cluster	H	M	L	Policy
C0	-	-	-	FedAvg-Random (Baseline)
C1	20	0	0	Performance
C2	15	5	0	
C3	10	5	5	
C4	5	10	5	
C5	5	5	10	
C6	0	5	15	
C7	0	0	20	Power

Table 5: Global parameter settings.

Setting	B	E	K
S1	32	10	20
S2	32	5	20
S3	16	5	20
S4	16	5	10

Based on the measured performance and power consumption, we evaluate the energy efficiency of participant devices in FL. To characterize the FL energy efficiency, we compare the energy efficiency of the various participant device clusters (Table 4) in Section 3. Based on the characterization results, we build AutoFL as described in Section 4, and implement it upon the FedAvg algorithm [64, 84] using PyTorch [97].

To evaluate the effectiveness of AutoFL, we compare it with five other design points:

- the FedAvg-Random baseline [84] where K participants are chosen randomly,
- Power where K participants are determined by minimizing power draw (i.e., C7 in Table 4),
- Performance where K participants are selected to achieve the best time performance (i.e., C1 in Table 4),
- $O_{participant}$ where the optimal cluster of K participants is determined by considering heterogeneity and runtime variance, and
- O_{FL} that considers available on-device co-processors for energy efficiency improvement over $O_{participant}$.

We also compare AutoFL with two closely related prior works: FedNova [121] and FEDL [26].

5.2 Workloads and Execution Scenarios

Workloads: We evaluate AutoFL using two common FL workloads: (1) training a CNN model with the MNIST dataset (**CNN-MNIST**) for image classification [68, 69, 113] and (2) training an LSTM model with the Shakespeare dataset (**LSTM-Shakespeare**) for the next character prediction [64, 84]. The workloads represent state-of-the-art FL use cases [28, 64, 73, 84]. In addition, we complement CNN-MNIST and LSTM-Shakespeare with another workload: (3) training the MobileNet model with the ImageNet dataset (**MobileNet-ImageNet**) for image classification [23, 47]. Table 5 summarizes the value ranges of the global parameters we consider in this work. Note, once the global parameters are determined for an FL use case, the values remain fixed until the model convergence [11, 64].

Runtime variance: To emulate realistic on-device interference, we initiate a synthetic co-running application on a random subset of devices, mimicking the effect of a real-world application (e.g., web browsing [49, 56, 95, 110, 111]). The synthetic application generates CPU and memory utilization patterns following the observed utilization patterns of web browsing. In addition, since the real-world network variability typically follows a Gaussian distribution [25], we emulate the random network bandwidth with a Gaussian distribution by adjusting the network delay.

Data distribution: We emulate different levels of data heterogeneity by distributing the total training dataset in four different ways [15, 75]: Ideal IID, Non-IID (50%), Non-IID (75%), and Non-IID (100%). In case of Ideal IID, all the data classes are evenly distributed to the devices in the cluster. On the other hand, in case of Non-IID (M%), M% of total devices have non-IID data while the rest have IID samples of all data classes. For non-IID devices, we distribute each data class randomly following a Dirichlet distribution with a 0.1 concentration parameters [15, 57, 72, 75, 78] — the smaller the value of the concentration parameter, the more each data class is concentrated on one device.

5.3 AutoFL Design Specification

Actions: We determine the 2-level actions for AutoFL. The first-level action determines a cluster of participant devices (Section 4.1) whereas the second-level action determines an execution target for the FL. Since the energy efficiency of local devices can be further improved via DVFS when stragglers are present, we identify V/F steps available in the FL system [5] as the augmented second-level action. Note, we measure the power consumption of different mobile devices at various frequency steps, in order to accurately model the energy efficiency of the FL execution.

Hyperparameters: There are two hyperparameters in the FL system: the learning rate and the discount factor. To determine them, we evaluate three values of 0.1, 0.5, and 0.9 for each one [20, 63]. We observe that the learning rate of 0.9 shows 20.1% and 32.5% better prediction accuracy than 0.5 and 0.1, respectively, meaning the more portion of the reward is added to the Q values, the better AutoFL works. This is because AutoFL needs to adapt to the runtime variance and data heterogeneity during the limited aggregation rounds. On the other hand, we observe that the discount factor of 0.1 shows 20.1% and 53.4% better prediction accuracy than 0.5 and 0.9, respectively, meaning the less portion of reward value for the next state is added to that for the current state, the better AutoFL works. This is because the consecutive states have a weak relationship owing to their stochastic nature, so that giving less weight to the reward in the near future improves the efficiency of AutoFL. Thus, in our evaluation, we use 0.9 and 0.1 for the learning rate and the discount factor, respectively.

6 EVALUATION RESULTS AND ANALYSIS

6.1 Result Overview

Compared with the baseline settings of FedAvg-Random, Power, and Performance, AutoFL improves the average FL energy efficiency of CNN-MNIST, LSTM-Shakespeare, and MobileNet-ImageNet by 4.3x, 3.2x, and 2.0x, respectively. It also exhibits better training




Figure 8: AutoFL improves the convergence time and the energy efficiency of FL, while also increasing model quality. It achieves 4.0x, 3.7x, and 5.1x higher energy efficiency than the baseline FedAvg-Random for CNN-MNIST, LSTM-Shakespeare, and MobileNet-ImageNet, respectively.

accuracy. Figure 8 compares the energy efficiency in performance-per-watt (PPW), the convergence time, and the training accuracy for the respective FL use cases, where PPW and the convergence time improvement are normalized to the FedAvg-Random baseline.

The energy efficiency gains of AutoFL come from two major sources. First, it can accurately identify near-optimal participants among a wide variety for each FL use case, reducing the performance slack of the stragglers. As a result, it improves the training time per round by an average of 3.5x, 2.9x, and 1.8x, over FedAvg-Random, Power, and Performance, respectively. This leads to faster convergence time. Second, AutoFL identifies more energy efficient execution targets for the individual participants. The energy efficiency is thus improved further by an average of 19.8% over *O_{participant}*. Compared with *O_{participant}*, AutoFL and *O_{FL}* experience slightly higher convergence time. This is because AutoFL leverages the remaining performance slack by considering alternative on-device execution targets and DVFS settings despite the slight increase in computation time.

In the case of CNN-MNIST and MobileNet-ImageNet, compute-intensive CONV and FC layers are dominant. In this case, performance-oriented selection (i.e., Performance) shows much higher energy efficiency than power-oriented selection (i.e., Power) due to the higher computation and memory capabilities of high-end devices. On the other hand, in the case of LSTM-Shakespeare, compute- and memory-intensive RC layers are dominant. In this case, the difference between the performance-oriented selection (i.e., Performance), and the power-efficient selection (i.e., Power) decreases. Nevertheless, since the baseline settings neglect to consider the NN characteristics explicitly in the participant device selection process, AutoFL's achieved energy efficiency outweighs that of other design points.

6.2 Adaptability and Accuracy Analysis

Adaptability to global parameters: AutoFL improves the energy efficiency and convergence time of various global parameter combinations. Figure 9 shows the average energy efficiency and convergence time of CNN-MNIST across four global parameter settings: S1 to S4 (Table 4 in Section 5.2). Although the optimal device cluster




Figure 9: Across the (B, E, K) global parameter settings of S1–S4, AutoFL achieves better training time performance and higher energy efficiency consistently.




Figure 10: In the presence of runtime variance, AutoFL can consistently and significantly improve the time-to-convergence and energy efficiency for FL in different execution environments.

varies along with the global parameters (as we observed in Section 3), AutoFL tries to accurately predict the near-optimal cluster of participant devices regardless of global parameters. Hence, it always outweighs the baseline settings of FedAvg-Random, Performance, and Power in terms of energy efficiency and convergence time. Moreover, since AutoFL also accurately predicts the near-optimal execution targets for individual devices, it achieves 15.9% better energy efficiency than $O_{participant}$.

Adaptability to stochastic variance: AutoFL can improve the energy efficiency and convergence time in the presence of stochastic on-device interference and network variance, independently. Figure 10 shows the PPW, convergence time, and training accuracy of CNN-MNIST, (a) when there is no on-device interference, (b) when there is on-device interference from co-running applications, and (c) when there is network variance. Even in the presence of runtime variance, AutoFL improves the average energy efficiency by 5.1x, 6.9x, and 2.6x, compared to FedAvg-Random, Power, and Performance. Note other NNs show similar result trends.

In the presence of runtime variance, the training time per round of the baseline settings significantly increases because of the increased on-device computation time or communication time. Even worse, since FedAvg excludes the severe stragglers from the round, the convergence time and the training accuracy is degraded as well. On the other hand, AutoFL selects the near-optimal cluster of




Figure 11: By explicitly taking into account data heterogeneity when selecting K devices, AutoFL achieves 4.0x, 5.5x, 9.3x, and 7.3x higher energy efficiency than the baseline FedAvg-Random for the four data distribution scenarios: (a)–(d).

participant devices even in the presence of runtime variance, mitigating the straggler problem. By doing so, it improves the convergence time by 3.4x, 3.3x, and 2.3x, compared with FedAvg-Random, Power, and Performance, respectively. Additionally, AutoFL also exploits the increased performance gap from the stragglers, improving 26.3% more energy efficiency compared to $O_{participant}$ at the cost of training time per round. As a result, AutoFL achieves similar energy efficiency, convergence time, and training accuracy with O_{FL} .

Adaptability to data heterogeneity: In the presence of data heterogeneity, AutoFL improves the energy efficiency by 7.4x, 5.5x, and 4.3x, respectively, compared with FedAvg-Random, Power, and Performance. It also exhibits much better convergence time and training accuracy. Figure 11 illustrates the energy efficiency, convergence time, and training accuracy of CNN-MNIST. Each column shows the varying level of data heterogeneity: (a) Ideal IID, (b) Non-IID (50%), (c) Non-IID (75%), (d) Non-IID (100%). Note other NNs also show similar result trends.

When there exist non-IID participants, the baseline settings (i.e., FedAvg-Random, Power, and Performance) that neglect to consider data heterogeneity experience sub-optimal energy efficiency, convergence time, and training accuracy. This is because naively including non-IID participants can significantly deteriorate model convergence — in the case of Non-IID (75%) and Non-IID (100%), CNN-MNIST with the baseline settings does not even converge in 1000 rounds (Figure 11(c) and Figure 11(d)). In contrast, AutoFL *learns the impact of data heterogeneity on the convergence time and energy efficiency dynamically and adapts to the different level of data heterogeneity among the devices*. Therefore, it achieves near-optimal energy efficiency, convergence time, and model quality even in the presence of data heterogeneity.

Prediction accuracy: AutoFL accurately selects the near-optimal participant cluster in varying data heterogeneity and runtime variance for the given NNs. Figure 12 shows how AutoFL and O_{FL} make the participant selection on three different categories of devices. For participant selection, AutoFL achieves 93.9% average prediction accuracy.

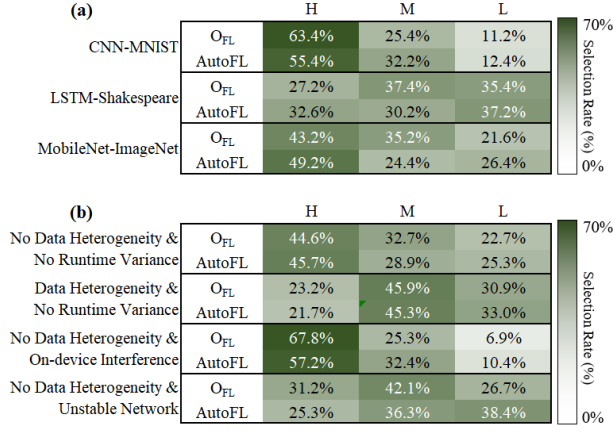


Figure 12: AutoFL can accurately track the decisions from the optimal policy (O_{FL}).

AutoFL selects the near-optimal participant cluster for different NNs. In Figure 12(a), the optimal selection varies with the NN characteristics. For example, O_{FL} includes more high-end devices for CNN-MNIST and MobileNet-ImageNet, whereas it includes more mid-end and low-end devices for LSTM-Shakespeare. AutoFL accurately captures those trends, achieving 94.2% average accuracy.

AutoFL also accurately adapts to the data heterogeneity and runtime variance. As shown in Figure 12(b), even in the presence of runtime variance and data heterogeneity, AutoFL makes the near-optimal participant selection, achieving 93.7% prediction accuracy on average.

AutoFL also selects the near-optimal execution targets in individual participant device. In the absence of runtime variance, CPU exhibits better energy efficiency than GPU, because of compute- and memory-intensive nature of training workloads. On the other hand, when on-device interference exists, the optimal execution target usually shifts from CPU to GPU, as the CPU performance is degraded due to 1) the competition for CPU time slices and cache, and 2) frequent thermal throttling. In case of unstable network, CPU and GPU show similar energy efficiency as FL becomes communication bound. AutoFL accurately captures such impact of runtime variance on optimal execution targets, achieving 92.9% average accuracy. Hence, as shown in Figure 8, 9, 10, and 11, AutoFL substantially improves energy efficiency, compared with $O_{participant}$, which does not select the optimal execution target in each participant.

6.3 Comparison with Prior Work

We compare AutoFL with two closely related prior works: FedNova [121] and FEDL [26]. FedNova normalizes gradient updates from stragglers or non-IID devices to those from ideal devices, whereas FEDL lets each client device to approximately adjust gradient updates based on the global weights. Both FedNova and FEDL allow partial updates from stragglers but implement random participant selections without considering runtime variance, such as on-device interference and network variability. Furthermore, neither exploits other available execution targets to improve FL performance or energy efficiency. On average, compared with FedNova and FEDL, AutoFL achieves 49.8% and 39.3% higher energy




Figure 13: Compared with FedNova [121] and FEDL [26], AutoFL achieves better convergence time and energy efficiency.

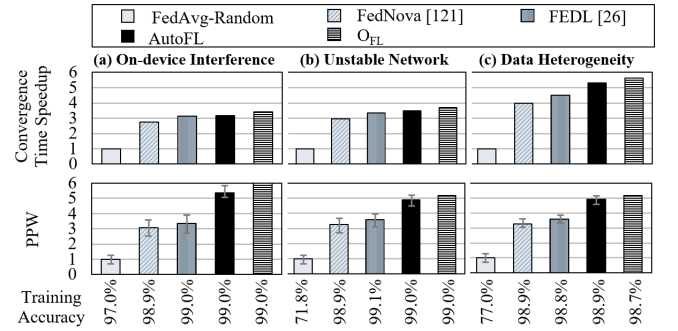


Figure 14: AutoFL outperforms both FedNova [121] and FEDL [26], even in the presence of runtime variance (a)(b) and data heterogeneity (c).

efficiency, respectively (Figure 13). In the presence of stochastic variance, FedNova and FEDL improve the execution time performance and PPW over the baseline, as expected. Similarly, AutoFL can further increase PPW by 62.7% and 48.8% over FedNova and FEDL, respectively (Figure 14).

Compared with the baseline, FedNova and FEDL are robust to data heterogeneity by giving less weights to gradient updates from non-IID devices. Nonetheless, including non-IID users can degrade model convergence, degrading time-to-convergence and energy efficiency. By contrast, AutoFL achieves near-optimal energy efficiency, convergence time, and model quality, even in the presence of data heterogeneity.

6.4 Overhead Analysis

Figure 15 shows that when training per-device Q-tables from scratch, the reward converges after about 50-80 aggregation rounds on average; more than 200 rounds are usually required for FL convergence. Before convergence, AutoFL exhibits 28.3% lower average energy efficiency than O_{FL} owing to the design space explorations. Nevertheless, it still achieves 52.1% energy saving against FedAvg-Random. After the reward is converged, AutoFL accurately selects the participants and execution targets, as we observed in Section 6.2. As a result, it can achieve 5.2x energy efficiency improvement for the entire FL, on average.




Figure 15: The reward is usually converged in 50-80 aggregations rounds. Sharing the Q-tables among the same category of devices expedites convergence.

The training overhead from the explorations can be alleviated by using the shared Q-tables. As shown in Figure 15, when the learned results are shared across the same category of devices, the RL training converges more rapidly, reducing the average training overhead by 29.3%; the prediction accuracy of AutoFL is slightly degraded by 2.7% though. This implies that, although each user experiences a different degree of runtime variance and data heterogeneity, learned results from various devices complement one another.

The runtime cost of training per-device Q-tables is 531.5 μ s, on average, excluding the time for FL execution. It corresponds to 0.8% of the average time for aggregation rounds. The overhead consists of observing the per-device states (496.8 μ s), selecting participants and execution targets based on the per-device Q-tables (10.5 μ s), calculating the reward (2.1 μ s), and updating the Q-tables (22.1 μ s). The overhead from training computation can be further alleviated by leveraging idle cores in mobile SoCs — the average thread-level parallelism for mobile applications is about 2 [31, 61], which is usually less than the number of available cores. Although AutoFL uses per-device Q-tables, its total memory requirement is feasible — for our experiments with 200 devices, the total requirement is 80MB, 0.25% of the typical 32GB DRAM capacity of commodity cloud servers. During the inference phase, misprediction contributed a negligible 5.6% timing and 8.8% energy efficiency overhead. AutoFL achieves an overall 93.8% prediction accuracy.

7 RELATED WORK

Energy optimization for mobile: Several prior works proposed statistical models to capture uncertainties in the mobile environment for dynamic energy management [32, 33, 110]. For example, Gaudette et al. proposed arbitrary polynomial chaos expansions as a way to consider the effect of uncertainties on mobile user experience [33]. Other computation offloading techniques also consider the performance variability of the mobile environment for energy efficiency optimization [2, 3, 22, 48, 60, 62, 63, 94, 101, 136, 139]. Although the aforementioned techniques addressed similar runtime variance in the edge-cloud execution environment, prior works are sub-optimal for FL because of the highly distributed nature of FL use cases—not only that system and data heterogeneity can easily degrade the quality of FL, but runtime variance can also introduce uncertainties in the training time and execution efficiency of federated learning.

Optimization for FL: FL enables a large cluster of decentralized mobile devices to collaboratively train a shared ML model while keeping the raw training samples on device [11, 34, 52, 64, 71, 77, 84, 116, 120, 135, 137]. To enable efficient FL deployment at the edge, FedAvg has been considered as the de facto FL algorithm [64, 84], maximizing the computation-communication ratio by using few participant devices with more per-device training iterations [73, 84, 112]. On top of FedAvg, various works have attempted to improve the model accuracy [28, 71, 74] or security robustness [34, 76, 79]. While FedAvg has opened up the possibility for practical FL deployment, it faces critical optimization challenges.

The high degree of system heterogeneity and stochastic edge-cloud execution environment introduces the straggler problem in FL, where the training time of each aggregation round is limited by the slowest device. To mitigate the straggler problem, previous works excluded stragglers from aggregation rounds [84, 87] or allowed asynchronous gradient updates [18, 119]. These approaches, however, often sacrifice accuracy, because of insufficient gradient updates. On the other hand, Zhan et al. tried to exploit the stragglers for power savings by adjusting the CPU frequency only [135]. However, their technique can increase the overall training time.

Varying characteristics of training samples per device introduce additional challenges to FL optimization. In particular, devices with non-IID data can degrade model quality and convergence time [15, 75]. To mitigate the effect of data heterogeneity, previous approaches proposed to exclude updates from non-IID devices with asynchronous aggregation algorithms [17, 18], to warm up the global model with a subset of globally shared data [141], or to share data across a subset of devices [28, 71]. However, none of the aforementioned techniques explicitly takes into account the stochastic runtime variance observed at the edge while simultaneously handling data and system heterogeneity. Another FedAvg-based algorithm, FedProx [73], handles system and data heterogeneity by allowing partial updates from stragglers and from devices with non-IID training data distributions. However, since FedProx applies the same partial update rate to the randomly selected participants, it still fails to address the heterogeneity and stochastic runtime variance that can come from those randomly selected participants. In practice, AutoFL can be used with FedProx for improving the device selection approach.

Finally, there has been very little work on FL energy efficiency optimization. Most prior work assume that FL is only activated when

smartphones are plugged into wall power [17, 99, 120, 130, 135] limiting its practicality. To the best of our knowledge, AutoFL is the first work that demonstrates the potential of energy-efficient FL execution in the presence of realistic in-the-field effects: system and data heterogeneity with performance uncertainties. By customizing a reinforcement learning-based approach, AutoFL can identify a near-optimal device cluster and respective execution targets, adapting to heterogeneity and runtime variance.

8 CONCLUSION

Federated learning has shown great promises in various applications with security-guarantee. To enable energy efficient FL on energy-constrained mobile devices, we propose a lightweight adaptive framework — AutoFL. The in-depth characterization of FL in edge-cloud systems demonstrates that an optimal cluster of participants and execution targets depend on various features: FL use cases, device and data heterogeneity, and runtime variance. AutoFL continuously learns and identifies a near-optimal cluster of participant devices and their respective execution targets by taking into account the aforementioned features. We design and construct representative FL use cases deployed in an emulated mobile cloud execution environment using off-the-shelf systems. On average, AutoFL improves FL energy efficiency by 5.2x, compared with the baseline setting of random selection, while improving convergence time and accuracy. By considering runtime variance along with system and data heterogeneity, AutoFL achieves an average of 49.8% and 39.3% higher energy efficiency, compared to the state-of-the-art techniques, FedNova [121] and FEDL [26], respectively. We demonstrate that AutoFL is a viable solution and will pave the path forward by enabling future work on energy efficiency improvement for FL in realistic execution environment.

ACKNOWLEDGMENTS

This work is supported in part by the National Science Foundation under grants CCF-1652132 and CCF-1618039 and by the National Research Foundation of Korea under grant NRF-2021R1C1C1008617.

REFERENCES

- [1] Bilge Acun, Matthew Murphy, Xiaodong Wang, Jade Nie, Carole-Jean Wu, and Kim Hazelwood. 2021. Understanding Training Efficiency of Deep Learning Recommendation Models at Scale. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*.
- [2] Taha Alfaqih, Mohammad Mehdi Hassan, Abdu Guma, Claudio Savaglio, and Giancarlo Fortino. 2020. Task Offloading and Resource Allocation for Mobile Edge Computing by Deep Reinforcement Learning Based on SARSA. *IEEE Access* 8 (2020), 54074–54084.
- [3] Majid Altamimi, Atef Abdrabou, Kshirasagar Naik, and Amiya Nayak. 2015. Energy Cost Models of Smartphones for Task Offloading to the Cloud. *IEEE Transactions on Emerging Topics in Computing* 3 (2015), Issue 3.
- [4] Amazon. 2021. Alexa. (2021). <https://developer.amazon.com/en-US/alexa>
- [5] Amazon. 2021. Amazon EC2. (2021). <https://aws.amazon.com/ec2>
- [6] Android. 2021. Android Neural Networks API. (2021). <https://developer.android.com/ndk/guides/neuralnetworks>
- [7] Apple. 2021. CoreML. (2021). <https://developer.apple.com/documentation/coreml>
- [8] Apple. 2021. Siri. (2021). <https://www.apple.com/siri>
- [9] Ammar Ahmad Awan, Hari Subramoni, and Dhabaleswar K. Panda. 2017. An In-depth Performance Characterization of CPU- and GPU-based DNN Training on Modern Architectures. In *Proceedings of the Machine Learning on HPC Environments (MLHPC)*.
- [10] Fabian Berns and Christian Biece. 2021. Complexity-Adaptive Gaussian Process Model Inference for Large-Scale Data. In *Proceedings of the SIAM International Conference on Data Mining (SDM)*.
- [11] Keith Bonawitz, Hubert Eichner, Wolfgang Grieskamp, Dmitriy Huba, Alex Ingerman, Vladimir Ivanov, Chloe Kiddon, Jakub Konecny, Stefano Mazzocchi, H B McMahan, Timon Van Overveldt, David Petrou, Daniel Ramage, and Jason Roselander. 2019. Towards Federated Learning at Scale: System Design. *arXiv:1902.01046* (2019).
- [12] Theodora S. Brisimi, Ruidi Chen, Theofanie Mela, Alex Olshevsky, Ioannis Ch Paschalidis, and Wei Shi. 2018. Federated Learning of Predictive Models from Federated Electronic Health Records. *International Journal of Medical Informatics* 112 (2018), 59–67.
- [13] David Brooks, Vivek Tiwari, and Margaret Martonosi. 2000. Wattch: A Framework for Architectural-level Power Analysis and Optimizations. In *Proceedings of the International Symposium on Computer Architecture (ISCA)*.
- [14] Ermao Cai, Da-Cheng Juan, Dimitrios Stamoulis, and Diana Maculescu. 2017. NeuralPower: Predict and Deploy Energy-Efficient Convolutional Neural Networks. In *Proceedings of the Asian Conference on Machine Learning (ACML)*.
- [15] Zheng Chai, Hannan Fayyaz, Zeshan Fayyaz, Ali Anwar, Yi Zhou, Heiko Ludwig, and Yue Cheng. 2019. Towards Taming the Resource and Data Heterogeneity in Federated Learning. In *Proceedings of the USENIX Conference on Operational Machine Learning (OpML)*.
- [16] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Linmin Zheng, Eddie Yan, Meghan Cowan, Haichen Shen, Leyuan Wang, Yuwei Hu, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. 2018. TVM: An Automated End-to-End Optimizing Compiler for Deep Learning. In *Proceedings of the USENIX Symposium on Operating Systems Design and Implementation*.
- [17] Yitao Chen, Saman Biokhaghazadeh, and Ming Zhao. 2019. Exploring the Capabilities of Mobile Devices in Supporting Deep Learning. In *Proceedings of the ACM/IEEE Symposium on Edge Computing (SEC)*. 127–138.
- [18] Yujing Chen, Yue Ning, Martin Slawski, and Huzefa Rangwala. 2019. Asynchronous Online Federated Learning for Edge Devices with Non-IID Data. *arXiv:1911.02134* (2019).
- [19] Yujing Chen, Yue Ning, Martin Slawski, and Huzefa Rangwala. 2020. Asynchronous Online Federated Learning for Edge Devices with Non-IID Data. In *Proceedings of IEEE International Conference on Big Data*.
- [20] Yonghun Choi, Seonghoon Park, and Hojung Cha. 2019. Optimizing Energy Efficiency of Browsers in Energy-aware Scheduling-enabled Mobile Devices. In *Proceedings of the International Conference on Mobile Computing and Networking (MobiCom)*.
- [21] David Cortes. 2018. Adapting Multi-Armed Bandits Policies to Contextual Bandits Scenarios. *arXiv:1811.04383* (2018).
- [22] Eduardo Cuervo, Aruna Balasubramanian, Dae ki Cho, Alec Wolman, Stefan Soriu, Ranveer Chandra, and Paramvir Bahl. 2010. MAUI: Making smartphone last longer with code offload. In *Proceedings of the International Conference on Mobile Systems, Applications, and Services (MobiSys)*.
- [23] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. ImageNet: A Large-Scale Hierarchical Image Database. In *Proceedings of the IEEE Computer Vision and Pattern Recognition (CVPR)*.
- [24] Ning Ding and Y Charlie Hu. 2017. GfxDoctor: A Holistic Graphics Energy Profiler for Mobile Devices. In *Proceedings of the European Conference on Computer Systems (EuroSys)*.
- [25] Ning Ding, Daniel Wagner, Xiaomeng Chen, Abhinav Pathak, Y Charlie Hu, and Andrew Rice. 2013. Characterizing and Modeling the Impact of Wireless Signal Strength on Smartphone Battery Drain. In *Proceedings of the International Conference on Measurement and Modeling of Computer Systems (SIGMETRICS)*. 29–40.
- [26] Canh T. Dinh, Nguyen H. Tran, Minh N. H. Nguyen, Choong Seon Hong, Wei Bao, Albert Y. Zomaya, and Vincent Gramoli. 2021. Federated Learning over Wireless Networks: Convergence Analysis and Resource Allocation. *IEEE/ACM Transactions on Networking* 29 (2021), 398–409. Issue 1.
- [27] DL4j. 2021. Deeplearning4j. (2021). <https://deeplearning4j.org>
- [28] Moming Duan, Xianzhang Chen, Renping Liu, Yujuan Tan, and Liang Liang. 2021. Self-Balancing Federated Learning with Global Imbalanced Data in Mobile Systems. *IEEE Transactions on Parallel and Distributed Systems* 32 (2021), 59–71. Issue 1.
- [29] Amir Erfan Eshratifar, Mohammad Saeed Abrishami, and Massoud Pedram. 2020. JointDNN: An Efficient Training and Inference Engine for Intelligent Mobile Cloud Computing Services. *IEEE Transactions on Mobile Computing* (2020).
- [30] Eyal Even-Dar, Shie Mannor, and Yishay Mansour. 2006. Action Elimination and Stopping Conditions for the Multi-Armed Bandit and Reinforcement Learning Problems. *Journal of Machine Learning Research* 7 (2006), 1079–1105.
- [31] Cao Gao, Anthony Gutierrez, Madhav Rajan, Ronald G. Dreslinski, Trevor Mudge, and Carole-Jean Wu. 2015. A Study of Mobile Device Utilization. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*.
- [32] Benjamin Gaudette, Carole-Jean Wu, and Sarma Vrudhula. 2016. Improving Smartphone User Experience by Balancing Performance and Energy with Probabilistic QoS Guarantee. In *Proceedings of the International Symposium on High Performance Computer Architecture*.

- [33] Benjamin Gaudette, Carole-Jean Wu, and Sarma Vrudhula. 2019. Optimizing User Satisfaction of Mobile Workloads Subject to Various Sources of Uncertainties. *IEEE Transactions on Mobile Computing* 18, 12 (2019), 2941–2953.
- [34] Robin C. Geyer, Tassilo Klein, and Moin Nabi. 2017. Differentially Private Federated Learning: A Client Level Perspective. *arXiv:1712.07557* (2017).
- [35] Zhangxiaowen Gong, Houxian Ji, Christopher Fletcher, Christopher Hughes, and Josep Torrellas. 2020. SparseTrain: Leveraging Dynamic Sparsity in Software for Training DNNs on General-Purpose SIMD Processors. In *Proceedings of the ACM International Conference on Parallel Architectures and Compilation Techniques (PACT)*.
- [36] Google. 2021. Google Cloud Vision. (2021). <https://cloud.google.com/vision>
- [37] Google. 2021. Google Pixel 5. (2021). https://store.google.com/us/product/pixel_5?hl=en-US
- [38] Google. 2021. Google Translate. (2021). <https://translate.google.com>
- [39] Google. 2021. Speech-to-Text. (2021). <https://cloud.google.com/speech-to-text>
- [40] Hui Guan, Laxmikant Kishor Mokadam, Xipeng Shen, Seung-Hwan Lim, and Robert Patton. 2020. FLEET: Flexible Efficient Ensemble Training for Heterogeneous Deep Neural Networks. In *Proceedings of Machine Learning and Systems (MLSys)*.
- [41] Myeonggyun Han, Jihoon Hyun, Seongbeom Park, Jinsu Park, and Woongki Baek. 2019. MOSAIC: Heterogeneity-, Communication-, and Constraint-aware Model Slicing and Execution for Accurate and Efficient Inference. In *Proceedings of the International Conference on Parallel Architecture and Compilation Techniques (PACT)*. 165–177.
- [42] Weituo Hao, Nikhil Mehta, Kevin J. Liang, Pengyu Cheng, Mostafa El-Khamy, and Lawrence Carin. 2020. WAFFLE: Weight Anonymized Factorization for Federated Learning. *arXiv:2008.05687* (2020).
- [43] Andrew Hard, Kanishka Rao, Rajiv Mathews, Swaroop Ramaswamy, Françoise Beaufays, Sean Augenstein, Hubert Eichner, Chloe Kiddon, and Daniel Ramage. 2018. Federated Learning for Mobile Keyboard Prediction. *arXiv:1811.03604* (2018).
- [44] Kim Hazelwood, Sarah Bird, David Brooks, Soumith Chintala, Utku Diril, Dmytro Dzhulgakov, Mohamed Fawzy, Bill Jia, Yangqing Jia, Aditya Kalro, James Law, Kevin Lee, Jason Lu, Pieter Noordhuis, Misha Smelyanskiy, Liang Xiong, and Xiaodong Wang. 2018. Applied Machine Learning at Facebook: A Datacenter Infrastructure Perspective. In *Proceedings of IEEE International Symposium on High-Performance Computer Architecture (HPCA)*.
- [45] Kaiming He, Georgia Gkioxari, Piotr Dollar, and Ross Girshick. 2018. Mask R-CNN. *arXiv:1703.06870v3* (2018).
- [46] Andrew Howard, Mark Sandler, Grace Chu, Liang-Chieh Chen, Bo Chen, Mingxing Tan, Weijun Wang, Yukun Zhu, Ruoming Pang, Vijay Vasudevan, Quoc V. Le, and Hartwig Adam. 2019. Searching for MobileNetV3. In *Proceedings of the International Conference on Computer Vision (ICCV)*.
- [47] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. 2017. MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications. *arXiv:1704.04861* (2017).
- [48] Liang Huang, Suzhi Bi, and Ying-Jun Angela Zhang. 2020. Deep Reinforcement Learning for Online Computation Offloading in Wireless Powered Mobile-Edge Computing Networks. *IEEE Transactions on Mobile Computing* (2020).
- [49] Yongbing Huang, Zhongbin Zha, Mingyu Chen, and Lixin Zhang. 2014. Moby: A Mobile Benchmark Suite for Architectural Simulators. In *Proceedings of the International Symposium on Performance Analysis of Systems and Software (ISPASS)*.
- [50] Huawei. 2021. Kirin 980, the World's First 7nm Process Mobile AI Chipset. (2021). <https://consumer.huawei.com/en/campaign/kirin980/>
- [51] Huawei. 2021. Kirin 990 Series, Rethink Evolution. (2021). <https://consumer.huawei.com/en/campaign/kirin-990-series/>
- [52] Sohei Itahara, Takayuki Nishio, Masahiro Morikura, and Koji Yamamoto. 2020. Lottery Hypothesis based Unsupervised Pre-training for Model Compression in Federated Learning. *arXiv:2004.09817* (2020).
- [53] Wenqi Jiang, Zhenhao He, Shuai Zhang, Thomas B. Preuber, Kai Zheng, Liang Feng, Jiansong Zhang, Tongxuan Liu, Yong Li, Jingren Zhou, and Ce Zhang. 2020. MicroRec: Efficient Recommendation Inference by Hardware and Data Structure Solutions. *arXiv:2010.05894* (2020).
- [54] Russ Joseph and Margaret Martonosi. 2001. Run-time Power Estimation in High Performance Microprocessors. In *Proceedings of International Symposium on Low Power Electronics and Design (ISLPED)*.
- [55] Norman P. Jouppi, Cliff Young, Nishant Patil, David Patterson, Gaurav Agrawal, Raminder Bajwa, Sarah Bates, Suresh Bhatia, Nan Boden, Al Borchers, Rick Boyle, Pierre-luc Cantin, Clifford Chao, Chris Clark, Jeremy Coriell, Mike Daley, Matt Dau, Jeffrey Dean, Ben Gelb, Tara Vazir Ghaemmaghami, Rajendra Gottipati, William Gulland, Robert Hagmann, C. Richard Ho, Doug Hogberg, John Hu, Robert Hundt, Dan Hurt, Julian Ibarz, Aaron Jaffey, Alek Jaworski, Alexander Kaplan, Harshit Khaitan, Daniel Killebrew, Andy Koch, Naveen Kumar, Steve Lacy, James Laudon, James Law, Diemthu Le, Chris Leary, Zhuyuan Liu, Kyle Lucke, Alan Lundin, Gordon MacKean, Adriana Maggiore, Maire Mahony, Kieran Miler, Rahul Nagarajan, Ravi Narayanaswami, Ray Ni, Kathy Nix, Thomas Norrie, Mark Omernick, Narayana Penukonda, Andy Phelps, Jonathan Ross, Matt Ross, Amir Salek, Emad Samadiani, Chris Severn, Gregory Sizikov, Matthew Snelham, Jed Souter, Dan Steinberg, Andy Swing, Mercedes Tan, Gregory Thorson, Bo Tian, Horia Toma, Erick Tuttle, Vijay Vasudevan, Richard Walter, Walter Wang, Eric Wilcox, and Doe Hyun Yoon. 2017. In-Datacenter Performance Analysis of a Tensor Processing Unit. In *Proceedings of IEEE/ACM International Symposium on Computer Architecture (ISCA)*.
- [56] Minh Ju, Hyeongyu Kim, and Soontae Kim. 2016. MofySim: A Mobile Full System Simulation Framework for Energy Consumption and Performance Analysis. In *Proceedings of International Symposium on Performance Analysis of Systems and Software (ISPASS)*.
- [57] Peter Kairouz, H. Brendan McMahan, Brendan Avent, Aurelien Bellet, Mehdi Bennis, Arjun Nitin Bhagoji, Kallista Bonawitz, Zachary Charles, Graham Cormode, Rachel Cummings, Rafael G. L. D'Oliveira, Hubert Eichner, Salim El Rouayheb, David Evans, Josh Gardner, Zachary Garrett, Adria Gascon, Badih Ghazi, Phillip B. Gibbons, Marco Gruteser, Zaid Harchaoui, Chaoyang He, Lie He, Zhouyuan Huo, Ben Hutchinson, Justin Hsu, Martin Jaggi, Tara Javidi, Gauri Joshi, Mikhail Khodak, Jakub Konecny, Aleksandra Korolova, Farinaz Koushanfar, Sanmi Koyejo, Tancrède Lepoint, Yang Liu, Prateek Mittal, Mehryar Mohri, Richard Nock, Ayfer Ozgur, Rasmus Pagh, Mariana Raykova, Hang Qi, Daniel Ramage, Ramesh Raskar, Dawn Song, Weikang Song, Sebastian U. Stich, Ziteng Sun, Ananda Theertha Suresh, Florian Tramer, Praneeth Vepakomma, Jianyu Wang, Li Xiong, Zheng Xu, Qiang Yang, Felix X. Yu, Han Yu, and Sen Zhao. 2019. Advances and Open Problems in Federated Learning. *arXiv:1912.04977* (2019).
- [58] Yiping Kang, Johann Hauswald, Cao Gao, Austin Rovinski, Trevor Mudge, Jason Mars, and Lingjia Tang. 2017. Neurosurgeon: Collaborative Intelligence between the Cloud and Mobile Edge. In *Proceedings of the International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. 615–629.
- [59] Sheng-Chun Kao, Geonhwa Jeong, and Tushar Krishna. 2020. ConfuciusX: Autonomous Hardware Resource Assignment for DNN Accelerators using Reinforcement Learning. In *Proceedings of the International Symposium on Microarchitecture (MICRO)*.
- [60] Young Geun Kim and Sung Woo Chung. 2017. Signal Strength-aware Adaptive Offloading for Energy Efficient Mobile Devices. In *Proceedings of the International Symposium on Low Power Electronics and Design*. 1–6.
- [61] Young Geun Kim, Minyong Kim, and Sung Woo Chung. 2017. Enhancing Energy Efficiency of Multimedia Applications in Heterogeneous Mobile Multi-core Processors. *IEEE Trans. Comput.* 66, 11 (2017), 1878–1889.
- [62] Young Geun Kim, Young Seo Lee, and Sung Woo Chung. 2020. Signal Strength-aware Adaptive Offloading with Local Image Preprocessing for Energy Efficient Mobile Devices. *IEEE Trans. Comput.* 69, 1 (2020), 99–101.
- [63] Young Geun Kim and Carole-Jean Wu. 2020. AutoScale: Energy Efficiency Optimization for Stochastic Edge Inference Using Reinforcement Learning. In *Proceedings of the IEEE/ACM International Symposium on Microarchitecture (MICRO)*. 1082–1096.
- [64] Jakub Konecny, H. Brendan McMahan, Felix X. Yu, Peter Richtarik, Ananda Theertha Suresh, and Dave Bacon. 2016. Federated Learning: Strategies for Improving Communication Efficiency. *arXiv:1610.05492* (2016).
- [65] D. E. Koulouriotis and A. Xanthopoulos. 2008. Reinforcement Learning and Evolutionary Algorithms for Non-stationary Multi-armed Bandit Problems. *Appl. Math. Comput.* 196 (2008).
- [66] Nicholas D. Lane, Sourav Bhattacharya, Petko Georgiev, Claudio Forlivesi, Lei Jiao, Lorena Qendro, and Fahim Kawsar. 2016. DeepX: A Software Accelerator for Low-power Deep Learning Inference on Mobile Devices. In *Proceedings of the International Conference on Information Processing in Sensor Networks (IPSN)*. 98–107.
- [67] Andrew Lavin and Scott Gray. 2016. Fast Algorithms for Convolutional Neural Networks. In *Proceedings of International Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [68] Yann LeCun, Leon Bottou, Yoshua Bengio, and Patrick Haffner. 1998. Gradient-based learning applied to document recognition. *Proc. IEEE* 86 (1998), 2278–2324. Issue 11.
- [69] Yann LeCun, Corinna Cortes, and Christopher J. C. Burges. 1998. The MNIST Database of handwritten digits. (1998). <http://yann.lecun.com/exdb/mnist/>
- [70] Dmitry Lepikhin, Hyukjoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. 2020. GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding. *arXiv:2006.16668* (2020).
- [71] Ang Li, Jingwei Sun, Binghui Wang, Lin Duan, Sicheng Li, Yiran Chen, and Hai Li. 2020. LotteryFL: Personalized and Communication-Efficient Federated Learning with Lottery Ticket Hypothesis on Non-IID Datasets. *arXiv:2008.03371* (2020).
- [72] Qibin Li, Yiqun Diao, Quan Chen, and Bingsheng He. 2021. Federated Learning on Non-IID Data Silos: An Experimental Study. *arXiv:2102.02079v2* (2021).
- [73] Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia Smith. 2020. Federated Optimization in Heterogeneous Networks. In *Proceedings of International Conference on Machine Learning and Systems*

- (MLSys).
- [74] Tian Li, Maziar Sanjabi, Ahmad Beirami, and Virginia Smith. 2020. Fair Resource Allocation in Federated Learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
 - [75] Xiang Li, Kaixuan Huang, Wenhao Yang, Shusen Wang, and Zhihua Zhang. 2020. On the Convergence of FedAvg on Non-IID Data. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
 - [76] Youjie Li, Mohammad Alian, Yifan Yuan, Zheng Qu, Peitian Pan, Ren Wang, Alexander Schwing, Hadi Esmaeilzadeh, and Nam Sung Kim. 2018. A Network-Centric Hardware/Algorithm Co-Design to Accelerate Distributed Training of Deep Neural Networks. In *Proceedings of the IEEE/ACM International Symposium on Microarchitecture (MICRO)*.
 - [77] Wei Yang Bryan Lim, Nguyen Cong Luong, Dinh Thai Hoang, Yutao Jiao, Ying-Chang Liang, Qiang Yang, Dusit Niyato, and Chunyan Miao. 2020. Federated Learning in Mobile Edge Networks: A Comprehensive Survey. *IEEE Communications Surveys & Tutorials* 22 (2020), 2031–2063. Issue 3.
 - [78] Tao Lin, Lingjing Kong, Sebastian U. Stich, and Martin Jaggi. 2020. Ensemble Distillation for Robust Model Fusion in Federated Learning. In *Proceedings of the International Conference on Neural Information Processing Systems (NIPS)*.
 - [79] Yunlong Lu, Xiaohong Huang, Yueyue Dai, Sabita Maharjan, and Yan Zhang. 2020. Differentially Private Asynchronous Federated Learning for Mobile Edge Computing in Urban Informatics. *IEEE Transactions on Industrial Informatics* 16 (2020), 2134–2143. Issue 3.
 - [80] Mi Luo, Fei Chen, Dapeng Hu, Yifan Zhang, Jian Liang, and Jiashi Feng. 2021. No Fear of Heterogeneity: Classifier Calibration for Federated Learning with Non-IID Data. *arXiv:2106.05001* (2021).
 - [81] Sumit K. Mandal, Ganapati Bhat, Janardhan Rao Doppa, Partha Pratim Pande, and Umit Y. Ogras. 2020. An Energy-aware Online Learning Framework for Resource Management in Heterogeneous Platforms. *ACM Transactions on Design Automation and Electronic Systems* (2020).
 - [82] Michael Mathieu, Mikael Henaff, and Yann LeCun. 2014. Fast Training of Convolutional Networks through FFTs. In *Proceedings of International Conference on Learning Representations (ICLR)*.
 - [83] Peter Mattson, Christine Cheng, Cody Coleman, Greg Diamos, Paulius Mickevicius, David Patterson, Hanlin Tang, Gu-Yeon Wei, Peter Bailis, Victor Bittorf, David Brooks, Dehao Chen, Debojyoti Dutta, Udit Gupta, Kim Hazelwood, Andrew Hock, Xinyuan Huang, Atsushi Ike, Bill Jia, Daniel Kang, David Kanter, Naveen Kumar, Jeffery Liao, Guokai Ma, Deepak Narayanan, Tayo Oguntebi, Gennady Pekhimenko, Lillian Pentecost, Vijay Janapa Reddi, Taylor Robie, Tom St. John, Carole-Jean Wu, Lingjie Xu, Masafumi Yamazaki, Cliff Young, and Matei Zaharia. 2020. MLPerf Training Benchmark. In *Proceedings of Machine Learning and Systems (MLSys)*.
 - [84] H. Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera-Arcas. 2017. Communication-Efficient Learning of Deep Networks from Decentralized Data. *arXiv:1602.05629* (2017).
 - [85] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fiedler, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharmashan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. 2015. Human-level Control through Deep Reinforcement Learning. *Nature* 518, 7540 (2015), 529–533.
 - [86] Sharada Mohanty, Jyotish Poonganam, Adrien Gaidon, Andrey Kolobov, Blake Wulfe, Dipam Chakraborty, Grazvydas Semetuskis, Joao Schapke, Jonas Kubilius, Jurgis Pasukonis, Linas Klimas, Matthew Hausknecht, Patrick MacAlpine, Quang Nhat Tran, Thomas Tumieli, Xiaocheng Tang, Xinwei Chen, Christopher Hesse, Jacob Hilton, William Hebggen Guss, Sahika Genc, John Schulman, and Karl Cobbe. 2021. Measuring Sample Efficiency and Generalization in Reinforcement Learning Benchmarks: NeurIPS 2020 Procgen Benchmark. *arXiv:2103.15332* (2021).
 - [87] Mehryar Mohri, Gary Sivek, and Ananda Theertha Suresh. 2019. Agnostic Federated Learning. In *Proceedings of International Conference on Machine Learning (ICML)*.
 - [88] Monsoon. 2021. High Voltage Power Monitor. (2021). <https://www.monsoon.com/high-voltage-power-monitor>
 - [89] Motorola. 2021. Moto X Force - Technical Specs. (2021). <https://support.motorola.com/uk/en/solution/MS112171>
 - [90] Maxim Naumov, John Kim, Dheevatsa Mudigere, Srinivas Sridharan, Xiaodong Wang, Whitney Zhao, Serhat Yilmaz, Changkyu Kim, Hector Yuen, Mustafa Ozdal, Krishnakumar Nair, Isabel Gao, Bor-Ying Su, Jiyan Yang, and Mikhail Smelyanskiy. 2020. Deep Learning Training in Facebook Data Centers: Design of Scale-up and Scale-out Systems. *arXiv:2003.09518* (2020).
 - [91] Rajiv Nishtala, Paul Carpenter, Vinicius Petrucci, and Xavier Martorell. 2017. Hipster: Hybrid Task Manager for Latency-Critical Cloud Workloads. In *Proceedings of the International Symposium on High Performance Computer Architecture (HPCA)*. 409–420.
 - [92] NVIDIA. 2021. NVIDIA TensorRT. (2021). <https://developer.nvidia.com/tensorrt>
 - [93] Santiago Pagani, Sai Manoj, Axel Jantsch, and Jorg Henkel. 2020. Machine Learning for Power, Energy, and Thermal Management on Multicore Processors: A Survey. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* 39, 1 (2020), 101–116.
 - [94] Shengli Pan, Zhiyong Zhang, Zongwang Zhang, and Deze Zeng. 2019. Dependency-Aware Computation Offloading in Mobile Edge Computing: A Reinforcement Learning Approach. *IEEE Access* 7 (2019), 134742–134753.
 - [95] Dhinakaran Pandiyan, Shin-Ying Lee, and Carole-Jean Wu. 2013. Performance, Energy Characterizations and Architectural Implications of an Emerging Mobile Platform Benchmark Suite. In *Proceedings of IEEE International Symposium on Workload Characterization (IISWC)*.
 - [96] Dhinakaran Pandiyan and Carole-Jean Wu. 2014. Quantifying the Energy Cost of Data Movement for Emerging Smart Phone Workloads on Mobile Platforms. In *Proceedings of IEEE International Symposium on Workload Characterization (IISWC)*.
 - [97] PyTorch. 2021. PyTorch. (2021). <https://pytorch.org>
 - [98] PyTorch. 2021. PyTorch Mobile. (2021). <https://pytorch.org/mobile/home/>
 - [99] Xinchu Qiu, Titouan Parcollet, Daniel J. Beutel, Taner Topal, Akhil Mathur, and Nicholas D. Lane. 2020. Can Federated Learning Save the Planet? *arXiv:2010.06537* (2020).
 - [100] Qualcomm. 2021. Qualcomm Neural Processing SDK for AI. (2021). <https://developer.qualcomm.com/software/qualcomm-neural-processing-sdk>
 - [101] Li Quan, Zhiliang Wang, and Fuji Ren. 2018. A Novel Two-Layered Reinforcement Learning for Task Offloading with Tradeoff between Physical Machine Utilization Rate and Delay. *Future Internet* 10 (2018), 1–17.
 - [102] Samyam Rajbhandari, Olatunji Ruwase, Jeff Rasley, Shaden Smith, and Yuxiong He. 2021. ZeRO-Infinity: Breaking the GPU Memory Wall for Extreme Scale Deep Learning. *arXiv:2104.07857* (2021).
 - [103] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, and Carole-Jean Wu. 2021. The Vision Behind MLPerf: Understanding AI Inference Performance. *IEEE Micro* 41 (2021), 10–18.
 - [104] Vijay Janapa Reddi, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, Maximilien Breughe, Mark Charlebois, William Chou, Ramesh Chukka, Cody Coleman, Sam Davis, Greg Diamos, Jared Duke, Dave Fick, J. Scott Gardner, Itay Hubara, Sachin Idgunji, Thomas B. Jablin, Jeff Jiao, Tom St. John, Pankaj Kanwar, David Lee, Jeffery Liao, Anton Lokhmotov, Francisco Massa, Peng Meng, Paulius Mickevicius, Colin Osborne, Gennady Pekhimenko, Arun Tejusve Raghunath Rajan, Dilip Sequeira, Ashish Sirasao, Fei Sun, Hanlin Tang, Michael Thomson, Frank Wei, Ephrem Wu, Lingjie Xu, Koichi Yamada, Bing Yu, George Yuan, Aaron Zhong, Peizhao Zhang, and Yuchen Zhou. 2020. MLPerf Inference Benchmark. In *Proceedings of ACM/IEEE Annual International Symposium on Computer Architecture (ISCA)*.
 - [105] Samsung. 2021. Exynos 9825 Processors. (2021). <https://www.samsung.com/semiconductor/minisite/exynos/products/mobileprocessor/exynos-9825/>
 - [106] Samsung. 2021. Exynos 990 Mobile Processors. (2021). <https://www.samsung.com/semiconductor/minisite/exynos/products/mobileprocessor/exynos-990/>
 - [107] Samsung. 2021. Samsung Galaxy S10e, S10, & S10+. (2021). <https://www.samsung.com/global/galaxy/galaxy-s10>
 - [108] Samsung. 2021. Samsung Neural SDK. (2021). <https://developer.samsung.com/neural/overview.html#Release-Notes>
 - [109] Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, and Liang-Chieh Chen. 2018. MobileNetV2: Inverted Residuals and Linear Bottlenecks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.
 - [110] Daves Shingari, Akhil Arunkumar, Benjamin Gaudette, Sarma Vrudhula, and Carole-Jean Wu. 2018. DORA: Optimizing Smartphone Energy Efficiency and Web Browser Performance under Interference. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 64–75.
 - [111] Daves Shingari, Akhil Arunkumar, and Carole-Jean Wu. 2015. Characterization and Throttling-based Mitigation of Memory Interference for Heterogeneous Smartphones. In *Proceedings of the IEEE International Symposium on Workload Characterization (IISWC)*.
 - [112] Virginia Smith, Chao-Kai Chiang, Maziar Sanjabi, and Ameet Talwalkar. 2017. Federated Multi-Task Learning. In *Proceedings of the International Conference on Neural Information Processing Systems (NIPS)*.
 - [113] Jost Tobias Springenberg, Alexey Dosovitskiy, Thomas Brox, and Martin Riedmiller. 2015. Striving for Simplicity: The All Convolutional Net. In *Proceedings of the International Conference on Language Representations (ICLR)*.
 - [114] Zhiqing Sun, Hongkun Yu, Xiaodan Song, Renjie Liu, Yiming Yang, and Denny Zhou. 2020. MobileBERT: A Compact Task-agnostic BERT for Resource-Limited Devices. *arXiv:2004.02984* (2020).
 - [115] Mingxing Tan and Quoc V. Le. 2019. EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *arXiv:1905.11946* (2019).
 - [116] Zeyi Tao and Qun Li. 2018. eSGD: Communication Efficient Distributed Deep Learning on the Edge. In *USENIX Workshop on Hot Topics in Edge Computing (HotEdge)*.
 - [117] TensorFlow. 2021. TFLite. (2021). <https://tensorflow.org/lite>
 - [118] Izzatul Umami and Lailia Rahmawati. 2021. Comparing Epsilon Greedy and Thompson Sampling Model for Multi-Armed Bandit Algorithm on Marketing

- Dataset. *Journal of Applied Data Sciences* 2 (2021), 14–26. Issue 2.
- [119] Marten van Dijk, Nhung V. Nguyen, Toan N. Nguyen, Lam M. Nguyen, Quoc Tran-Dinh, and Phuong Ha Nguyen. 2020. Asynchronous Federated Learning with Reduced Number of Rounds and with Differential Privacy from Less Aggregated Gaussian Noise. (2020).
 - [120] Cong Wang, Yuanyuan Yang, and Pengzhan Zhou. 2021. Towards Efficient Scheduling of Federated Mobile Devices under Computational and Statistical Heterogeneity. *IEEE Transactions on Parallel and Distributed Systems* 32 (2021), 394–410. Issue 2.
 - [121] Jianyu Wang, Qinghua Liu, Hao Liang, Gauri Joshi, and H. Vincent Poor. 2020. Tackling the Objective Inconsistency Problem in Heterogeneous Federated Optimization. In *Proceedings of Conference on Neural Information Processing Systems (NeurIPS)*.
 - [122] Siqi Wang, Anju Pathania, and Tulika Mitra. 2020. Neural Network Inference on Mobile SoCs. *IEEE Design & Test* (2020).
 - [123] Siqi Wang, Anju Pathania, and Tulika Mitra. 2020. Neural Network Inference on Mobile SoCs. *IEEE Design & Test* (2020).
 - [124] Yu Emma Wang, Gu-Yeon Wei, and David Brooks. 2020. A Systematic Methodology for Analysis of Deep Learning Hardware and Software Platforms. In *Proceedings of Machine Learning Systems (MLSys)*.
 - [125] Yu Emma Wang, Carole-Jean Wu, Xiaodong Wang, Kim Hazelwood, and David Brooks. 2020. Exploiting Parallelism Opportunities with Deep Learning Frameworks. *ACM Transactions on Architecture and Code Optimization* 18 (2020), 1–23. Issue 1.
 - [126] Bichen Wu, Xiaoliang Dai, Peizhao Zhang, Yanghan Wang, Fei Sun, Yiming Wu, Yuandong Tian, Peter Vajda, Yangqing Jia, and Kurt Keutzer. 2018. FBNet: Hardware-Aware Efficient ConvNet Design via Differentiable Neural Architecture Search. *arXiv:1812.03443* (2018). [arXiv:cs.CV/1812.03443](https://arxiv.org/abs/1812.03443)
 - [127] Bichen Wu, Xiaoliang Dai, Peizhao Zhang, Yanghan Wang, Fei Sun, Yiming Wu, Yuandong Tian, Peter Vajda, Yangqing Jia, and Kurt Keutzer. 2019. FBNet: Hardware-Aware Efficient ConvNet Design via Differentiable Neural Architecture Search. In *Proceedings of the International Conference on Computer Vision and Pattern Recognition*.
 - [128] Carole-Jean Wu, David Brooks, Kevin Chen, Douglas Chen, Sy Choudhury, Marat Dukhan, Kim Hazelwood, Eldad Isaac, Yangqing Jia, Bill Jia, Tommer Leyvand, Hao Lu, Yang Lu, Lin Qiao, Brandon Reagen, Joe Spisak, Fei Sun, Andrew Tulloch, Peter Vajda, Xiaodong Wang, Yanghan Wang, Bram Wasti, Yiming Wu, Ran Xian, Sungjoo Yoo, and Peizhao Zhang. 2019. Machine Learning at Facebook: Understanding Inference at the Edge. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 331–344.
 - [129] Carole-Jean Wu and Margaret Martonosi. 2011. Characterization and dynamic mitigation of intra-application cache interference. In *Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software*. 2–11.
 - [130] Zichen Xu, Li Li, and Wenting Zou. 2019. Exploring Federated Learning on Battery-powered Devices. In *Proceedings of the ACM Turing Celebration Conference*.
 - [131] Feng Yan, Olatunji Ruwase, Yuxiong He, and Trishul Chilimbi. 2015. Performance Modeling and Scalability Optimization of Distributed Deep Learning Systems. In *Proceedings of ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*.
 - [132] Miao Yang, Akitanoshou Wong, Hongbin Zhu, Haifeng Wang, and Hua Qian. 2020. Federated Learning with Class Imbalance Reduction. *arXiv:2011.11266* (2020).
 - [133] Denis Yarats, Amy Zhang, Ilya Kostrikov, Brandon Amos, Joelle Pineau, and Rob Fergus. 2019. Improving Sample Efficiency in Model-Free Reinforcement Learning from Images. *arXiv:1910.01741v3* (2019).
 - [134] Chunxing Yin, Blige Acun, Xing Liu, and Carole-Jean Wu. 2021. TT-Rec: Tensor Train Compression for Deep Learning Recommendation Model Embeddings. In *Proceedings of Machine Learning and Systems (MLSys)*.
 - [135] Yufeng Zhan, Peng Li, and Song Guo. 2020. Experience-Driven Computational Resource Allocation of Federated Learning by Deep Reinforcement Learning. In *Proceedings of the IEEE International Parallel and Distributed Processing Symposium (IPDPS)*.
 - [136] Bingxin Zhang, Guopeng Zhang, Weice Sun, and Kun Yang. 2020. Task Offloading with Power Control for Mobile Edge Computing Using Reinforcement Learning-based Markov Decision Process. *Mobile Information Systems* (2020).
 - [137] Jiale Zhang, Yanchao Zhao, Junyu Wang, and Bing Chen. 2020. FedMEC: Improving Efficiency of Differentially Private Federated Learning via Mobile Edge Computing. *Mobile Networks and Applications* 25 (2020), 2421–2433.
 - [138] Lide Zhang, Birjodh Tiwana, Zhiyun Qian, Zhaoguang Wang, Robert P Dick, Z Morley Mao, and Lei Yang. 2010. Accurate Online Power Estimation and Automatic Battery Behavior based Power Model Generation for Smartphones. In *Proceedings of the International Conference on Hardware/Software Codesign and System Synthesis*. 105–114.
 - [139] Tianyu Zhang, Yi-Han Chiang, Cristian Borcea, and Yusheng Ji. 2019. Learning-Based Offloading of Tasks with Diverse Delay Sensitivities for Mobile Edge Computing. In *Proceedings of the IEEE Global Communications Conference (GLOBECOM)*.
 - [140] Zhengming Zhang, Yaoqing Yang, Zhewei Yao, Yujun Yan, Joseph E. Gonzalez, and Michael W. Mahoney. 2020. Improving Semi-supervised Federated Learning by Reducing the Gradient Diversity of Models. *arXiv:2008.11364* (2020).
 - [141] Yue Zhao, Meng Li, Liangzhen Lai, Naveen Suda, Damon Civin, and Vikas Chandra. 2018. Federated Learning with Non-IID Data. *arXiv:1806.00582* (2018).
 - [142] Bojian Zheng, Abhishek Tiwari, Nandita Vijaykumar, and Gennady Pekhimenko. 2020. Echo: Compiler-based GPU Memory Footprint Reduction for LSTM RNN Training. In *Proceedings of the ACM/IEEE Annual Symposium on Computer Architecture (ISCA)*. 1089–1102.
 - [143] Guanwen Zhong, Akshat Dubey, Cheng Tan, and Tulika Mitra. 2019. Synergy: An HW/SW Framework for High Throughput CNNs on Embedded Heterogeneous SoC. *ACM Transactions on Embedded Computing Systems* 18, 2 (2019), 1–23.
 - [144] Yuhao Zhu and Vijay Janapa Reddi. 2013. High-performance and energy-efficient mobile web browsing on big/little systems. In *Proceedings of the IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 13–24.
 - [145] Barret Zoph, Vijay Vasudevan, Jonathon Shlens, and Quoc V Le. 2018. Learning Transferable Architectures for Scalable Image Recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 8697–8710.